






THE NEURAL SIMULATION TOOL NEST

CNS*2023, Leipzig

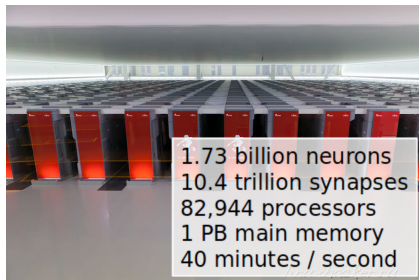
July 15th, 2023 | Charl Linssen (c.linssen@fz-juelich.de) | SimLab Neuroscience

WHEN TO USE NEST?



NEST = NEURAL SIMULATION TOOL

- Point neurons and neurons with few electrical compartments
 - Phenomenological synapse models (STDP, STP)
 - + gap junctions, neuromodulation and structural plasticity
 - Frameworks for rate models and binary neurons
 - Support for neuroscience interfaces (MUSIC, libneurosim)
-
- Highly efficient C++ core with a Python frontend
 - Hybrid parallelization (OpenMP+MPI)
 - Same code from laptops to supercomputers



NEST DESIGN GOALS

NEST development is always driven by scientific needs.

High accuracy and flexibility



- Exact integration is used for suitable neuron models
- Spike interaction in continuous time available for suitable neuron models
- Extremely scalable: same code from laptop to supercomputers

Constant quality assurance

- Automated unit test suite included in NEST build
- Continuous integration for all repository checkins
- Peer review for all code contributions

NEURONAL SIMULATIONS IN NEST

A simulation in NEST mimics a neuroscientific experiment.



Except that you can easily measure and control every neuron and connection in the network!

THREE MAIN COMPONENTS OF A NEST SIMULATION

■ Nodes

- Neurons – Devices (– Sub-networks)
- Have dynamic state variable(s) that changes over time ($V_m(t)$)
- Can be affected by events (spikes)

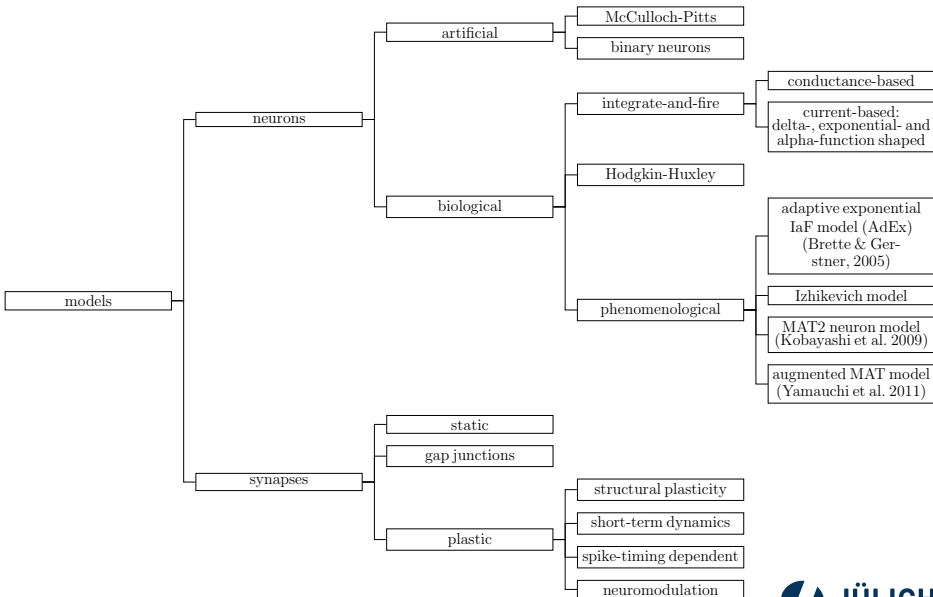
■ Events

- Pieces of information of a particular type (e.g., spike, voltage or current event)
- Recording devices: ‘spike_detector’, ‘voltmeter’, ‘multimeter’

■ Connections

- Communication channels for the exchange of events
- Directed (from source node to target node)
- Weighted (how strongly does an event influence the target node)
- Delayed (length of transmission duration between source and target)
- Connections are created using one global Connect function

NEST NEURON AND SYNAPSE MODELS



NEURON MODELS

- Integrate-and-fire models (iaf_)
 - Current-based (iaf_psc)
 - Conductance-based (iaf_cond)
 - Different post-synaptic shapes (_alpha, _exp, _delta)
- Single compartment Hodgkin-Huxley models (hh_)
- Adaptive exponential integrate-and-fire models (aeif_)
- Allen Institute generalised leaky integrate-and fire family (gif_)
- Hill-Tononi model (Hill & Tononi 2005)
- Neuron models with >1 compartments (“mesocompartmental”)
- ...and many more!

SYNAPSE MODELS

- Gap junctions (electrical synapse)
- Short term depression, facilitation
- Spike-timing dependent plasticity (STDP)
 - All-to-all; several nearest-neighbour variants
- Triplet STDP (Pfister & Gerstner 2006)
- Voltage-based STDP (Clopath et al. 2010)
- Dopamine-modulated STDP (Potjans et al. 2010)
- ...and many more!

STIMULATION DEVICES

Spike generators:

- `spike_generator` spikes at prescribed points in time
- `poisson_generator` spikes according to a Poisson distribution
- `gamma_sup_generator` spikes according to a Gamma distribution

Current generators

- `ac_generator` provides a sine-shaped current
- `dc_generator` provides a constant current
- `step_current_generator` provides a step-wise constant current
- `noise_generator` provides a random noise current


RECORDING DEVICES

- spike_detector records incoming spikes
 - multimeter records analog quantities (potentials, conductances, ...)
 - voltmeter records the membrane potential
 - correlation_detector records pairwise cross-correlations between the spiking activity of neurons
 - weight_recorder records the weight of connections
- can record to memory, file, or stream data!


SPECIFICATION OF CONNECTIVITY

- One-to-one
- All-to-all
- Fixed in-, out-degree
- Fixed total number (N)
- Pairwise Bernoulli (p)

```
A = Create('iaf_psc_alpha', n)  
B = Create('spike_detector', n)  
Connect(A, B, 'one_to_one')
```



```
A = Create("iaf_psc_alpha", n)  
B = Create("iaf_psc_alpha", m)  
Connect(A, B, {'rule': 'fixed_indegree',  
                'indegree': N})
```



RANDOMIZATION OF SYNAPSE PROPERTIES

Dozens of standard distributions are natively supported.


```
delay_dist = {'distribution': 'uniform',  
              'low': 0.8, 'high': 2.5}
```

```
alpha_dist = {'distribution': 'normal_clipped',  
              'low': 0.5, 'mu': 5.0,  
              'sigma': 1.0}
```


```
syn_dict = {'model': 'stdp_synapse',  
            'weight': 2.5,  
            'delay': delay_dist,  
            'alpha': alpha_dist}
```

STRUCTURED NETWORKS USING TOPOLOGY

- Set node positions on grids or arbitrary points in space (1D, 2D, 3D)
- Nodes can be neurons or combinations of neurons and devices
- Connect nodes in a position- and distance-dependent manner
- Set boundary condition (periodic or not)



GAP JUNCTIONS




Neuron i (hh_psc_alpha_gap)

$y'_i(t) = f_i(y_i(t))$, $y_i(t_0)$ given


$$\begin{aligned} \frac{V'_i}{C_m} = & -I_i^{ionic}(V_i, m_i, h_i, n_i, p_i) \\ & + I_i^{applied}(I_i^{ex}, I_i^{in}) \\ & + I_i^{gap}(V_i, V_j) \end{aligned}$$

- at each time point neuron i needs membrane potential of neuron j
- large system of differential equations
- naïve: communication of V in each step
- better: Jacobi waveform relaxation



Hahne et al. (2015). A unified framework for spiking and gap-junction interactions in distributed neural network simulations. Frontiers in Neuroinformatics. 9:22

NEST PERFORMANCE



Maximum network size and corresponding run time as function of number of virtual processes on the K computer (red) and JUQUEEN (blue). Taken from Kunkel et al., (2014), [Front Neuroinf](#). DOI: 10.3389/fninf.2014.00078

HOW TO USE NEST?

- The Python interface PyNEST

```
import nest
```

```
n = nest.Create("iaf_psc_exp", 8000, {"V_m": -65.})
```

```
sd = nest.Create("spike_detector")
```

```
nest.Connect(n, sd)
```

```
...
```

```
nest.Simulate(1000.)
```

- The multi-simulator Python interface PyNN

```
from pyNN import nest
```

```
nest.setup()
```

```
neuron_t = nest.native_cell_type("iaf_psc_exp")
```

```
pop_exc = nest.Population(8000, neuron_t, {})
```



```
pop_exc.initialize(V_m=-65.)
```

```
...
```

```
nest.run(1000.)
```

NEURONAL SIMULATIONS IN NEST

A simulation in NEST mimics a neuroscientific experiment.



Except that you can easily measure and control every neuron and connection in the network!

A FULL EXAMPLE


```
import nest                                     # import NEST module
neuron = nest.Create('iaf_psc_exp')            # create a neuron
voltmeter = nest.Create('voltmeter')          # create a voltmeter
spikegenerator = nest.Create('spike_generator') # create a spike generator
nest.SetStatus(spikegenerator, {'spike_times': [10., 50.]}) # let it spike

# connect spike generator and voltmeter to the neuron
nest.Connect(spikegenerator, neuron, syn_spec={'weight' : 1E3})
nest.Connect(voltmeter, neuron)

nest.Simulate(100.) # run the simulation

# read out recording time and voltage from voltmeter and plot them
times = nest.GetStatus(voltmeter)[0]['events']['times']
voltage = nest.GetStatus(voltmeter)[0]['events']['V_m']
pl.plot(times, voltage)
pl.xlabel('time (ms)'); pl.ylabel('membrane potential (mV)')
pl.show()
```

A FULL EXAMPLE



EVENT-DRIVEN VS. TIME-DRIVEN


Event-driven simulation:

- Visit a neuron only when it receives an event (e.g. a spike)
- From $y(t_i)$, calculate $y(t_{i+1})$

Time-driven simulation:


- Visit each neuron in each time step h
- From $y(ih)$, calculate $y([i + 1]h)$

EVENT-DRIVEN VS. TIME-DRIVEN



Hanuschkin et al. Front. Neuroinf. 2010

EVENT-DRIVEN VS. TIME-DRIVEN



Hanuschkin et al. Front. Neuroinf. 2010

EVENT-DRIVEN VS. TIME-DRIVEN

	Event-driven	Time-driven
Pros	<ul style="list-style-type: none">■ more efficient for low input rates■ 'correct' solution for invertible neuron models	<ul style="list-style-type: none">■ more efficient for high input rates■ works for all neuron models■ scales well
Cons	<ul style="list-style-type: none">■ only works for neurons with invertible dynamics■ event queue does not scale well	<ul style="list-style-type: none">■ only 'approximate' solution even for analytically solvable models■ spikes can be missed due to discrete sampling of membrane potential

EVENT-DRIVEN VS. TIME-DRIVEN

NEST uses a hybrid approach to simulation

- input events to neurons are frequent: time-driven algorithm
 - If the dynamics is nonlinear, we need a numerical method to solve it, e.g.:
 - Forward Euler: $y([i + 1]h) = y(ih) + h \cdot \dot{y}(ih)$
 - Runge-Kutta (k -th order)
 - Runge-Kutte-Fehlberg with adaptive step size
 - ...
- Use a pre-implemented solver, for example, from the GNU Scientific Library (GSL).
- If the dynamics is linear (e.g. leaky integrate-and-fire), we can solve it exactly.


EVENT-DRIVEN VS. TIME-DRIVEN

NEST uses a hybrid approach to simulation

- input events to neurons are frequent: time-driven algorithm
 - If the dynamics is nonlinear, we need a numerical method to solve it, e.g.:
 - Forward Euler: $y([i + 1]h) = y(ih) + h \cdot \dot{y}(ih)$
 - Runge-Kutta (k -th order)
 - Runge-Kutta-Fehlberg with adaptive step size
 - ...
- Use a pre-implemented solver, for example, from the GNU Scientific Library (GSL).
- If the dynamics is linear (e.g. leaky integrate-and-fire), we can solve it exactly.
- events at synapses are rare: event driven component
 - Exception: gap junctions

MINIMUM SYNAPTIC DELAY


During an interval of the minimal transmission delay in the network (Δ), neurons are effectively decoupled.



Input events to neurons are frequent: time-driven updates.
Events at synapses are rare: event driven.


COMMUNICATION OF EVENTS

- communication only required in intervals of the minimal delay between neurons




COMMUNICATION OF EVENTS

- communication only required in intervals of the minimal delay between neurons
- communication frequency independent of step size h




COMMUNICATION OF EVENTS

- communication only required in intervals of the minimal delay between neurons
- communication frequency independent of step size h
- less communications containing more data is more efficient due to overhead of communication between machines




REPRESENTATION OF NETWORK STRUCTURE: SERIAL

- Each neuron and synapse maintains its own parameters
- Synapses save the index of the target neuron




REPRESENTATION OF NETWORK STRUCTURE: DISTRIBUTED

- neurons are distributed round robin onto processes
- one target list for every neuron on each machine
- synapse stored on machine that hosts the target neuron
- wiring is a parallel operation



NESTML

NESTML is a domain-specific language for neuron and synapse models.



nest::ml

```
synapse stdp:
  state:
    w real = 1
    tr_post real = 0
    tr_pre real = 0

  equations:
    tr_pre' = -tr_pre / tau_tr
    tr_post' = -tr_post / tau_tr

  input:
    pre_spikes real <- spike
    post_spikes real <- spike

  onReceive(pre_spikes):
    w -= alpha * tr_post      # depress synapse
    tr_pre += 1              # update presynaptic trace
    deliver_spike(w, delay)  # to postsynaptic partner

  onReceive(post_spikes):
    w += alpha * tr_pre      # potentiate synapse
    tr_post += 1             # update postsynaptic trace

  parameters:
    delay ms = 1 ms         # dendritic delay
    tau_tr ms = 50 ms        # pre/post trace time const.
    alpha real = .02         # learning rate
```

NESTML: DESIGN PRINCIPLES


- Concise; low on boilerplate
- Speak in the vernacular of the neuroscientist (keywords such as neuron, synapse)
- Easy (dynamical) equation handling coupled with imperative-style programming (if $V_m \geq \text{threshold}$: ...)

NESTML comes with a code generation toolbox.

- Code generation (model definition but not instantiation)
- Automated ODE analysis and solver selection
- Flexible addition of targets using Jinja2 templates

THE MICROCIRCUIT MODEL


- 10^5 identical leaky-integrate and fire neurons
- $3 \cdot 10^8$ exponentially decaying synaptic currents
- four layers with one excitatory and one inhibitory population each
- size of populations and connection probabilities deduced from anatomical data sets
- asynchronous irregular and cell-type specific firing rates
- thalamic stimulation elicits flow of activity through cortical layers



Potjans and Diesmann (2014) The Cell-Type Specific Cortical Microcircuit: Relating Structure and Activity in a Full-Scale Spiking Network Model. *Cerebral Cortex* 24(3):785-806

THE MULTI-AREA MODEL

- full-density model of macaque visual cortex
- axonal tracing data from the CoCoMac database
- stable asynchronous irregular ground state
- produces realistic spiking statistics in V1
- functional connectivity compares to fMRI measurements

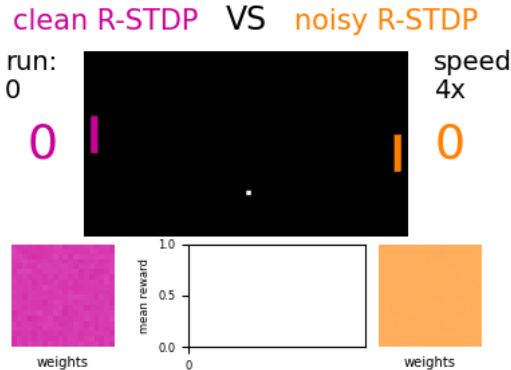


Schmidt et al. (2018) Multi-scale account of the network structure of macaque visual cortex. *Brain Structure and Function* 223(3):1409-1435

Schmidt et al. (2018) A multi-scale layer-resolved spiking network model of resting-state dynamics in macaque visual cortical areas. *PLOS CB* 14(10):e1006359

LEARNING TO PLAY PONG

Two spiking neural networks of two layers each compete by encoding an input-output mapping in their weights using reward-modulated STDP



Wunderlich T., et al (2019). Demonstrating advantages of neuromorphic computation: a pilot study. *Frontiers in neuroscience*, 13, 260

GETTING HELP

Within Python:

```
nest.help('iaf_psc_exp')  
nest.help('Connect')
```

Online documentation:

<https://nest-simulator.readthedocs.io/>

Community:

- NEST user mailing list
- Bi-weekly open video conference
- <http://github.com/nest/nest-simulator/>
- Annual NEST Conference: a forum for users and developers

Please tell us about problems. We can only fix what we know of!

REFERENCES AND FURTHER READING

- The NEST Simulator documentation at <https://nest-simulator.readthedocs.io/>
- Gewaltig et al. (2012) NEST by example: An introduction to the neural simulation tool NEST. doi:10.1007/978-94-007-3858-4_18
- Hanuschkin et al. (2010) A general and efficient method for incorporating precise spike times in globally time-driven simulations. doi:10.3389/fninf.2010.00113
- Jordan et al. (2018) Extremely Scalable Spiking Neuronal Network Simulation Code: From Laptops to Exascale Computers. doi:10.3389/fninf.2018.00002

ACKNOWLEDGMENTS

This presentation is based on previous work by many people.

- Hannah Bos
- David Dahmen
- Moritz Deger
- Jochen Martin Eppler
- Espen Hagen
- Abigail Morrison
- Jannis Schuecker
- Johanna Senk
- Tom Tetzlaff
- Sacha van Albada
- Charl Linssen
- Pooja Babu