




# THE NEURAL SIMULATION TOOL NEST

## CNS2020 full-day tutorial

July 18th, 2020 | Charl Linssen (c.linssen@fz-juelich.de) | SimLab Neuroscience

# WHEN TO USE NEST?




Possibility to simulate large networks

Complexity of single elements

# NEST = NEURAL SIMULATION TOOL

- Point neurons and neurons with few electrical compartments
  - Phenomenological synapse models (STDP, STP)
    - + gap junctions, neuromodulation and structural plasticity
  - Frameworks for rate models and binary neurons
  - Support for neuroscience interfaces (MUSIC, libneurosim)
- 
- Highly efficient C++ core with a Python frontend
  - Hybrid parallelization (OpenMP+MPI)
  - Same code from laptops to supercomputers



# NEST DESIGN GOALS

NEST development is always driven by scientific needs.

High accuracy and flexibility

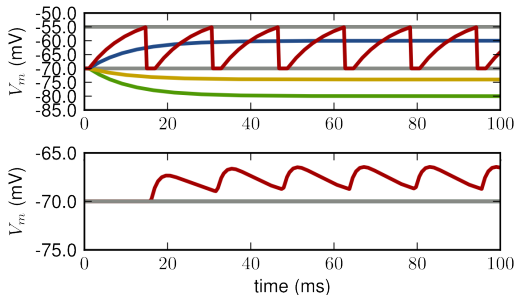
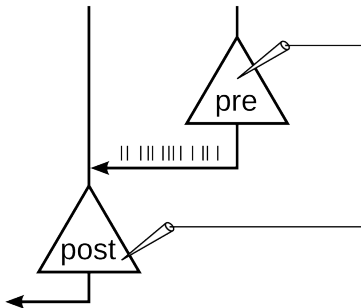
- Exact integration is used for suitable neuron models
- Spike interaction in continuous time available for suitable neuron models
- Extremely scalable: same code from laptop to supercomputers

Constant quality assurance

- Automated unit test suite included in NEST build
- Continuous integration for all repository checkins
- Peer review for all code contributions

# NEURONAL SIMULATIONS IN NEST

A simulation in NEST mimics a neuroscientific experiment.



# NEURON MODELS

- Integrate-and-fire models (iaf\_)
  - Current-based (iaf\_psc)
  - Conductance-based (iaf\_cond)
  - Different post-synaptic shapes (\_alpha, \_exp, \_delta)
- Single compartment Hodgkin-Huxley models (hh\_)
- Adaptive exponential integrate-and-fire models (aeif\_)
- Allen Institute generalised leaky integrate-and fire family (gif\_)
- Hill-Tononi model (Hill & Tononi 2005)
- Neuron models with  $>1$  compartments (“mesocompartmental”)
- ...and many more!

# SYNAPSE MODELS

- Gap junctions (electrical synapse)
- Short term depression, facilitation
- Spike-timing dependent plasticity (STDP)
  - All-to-all; several nearest-neighbour variants
- Triplet STDP (Pfister & Gerstner 2006)
- Voltage-based STDP (Clopath et al. 2010)
- Dopamine-modulated STDP (Potjans et al. 2010)
- ...and many more!

# STIMULATION DEVICES

## Spike generators:

- spike\_generator spikes at prescribed points in time
- poisson\_generator spikes according to a Poisson distribution
- gamma\_sup\_generator spikes according to a Gamma distribution

## Current generators

- ac\_generator provides a sine-shaped current
- dc\_generator provides a constant current
- step\_current\_generator provides a step-wise constant current
- noise\_generator provides a random noise current



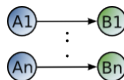
# RECORDING DEVICES

- `spike_detector` records incoming spikes
- `multimeter` records analog quantities (potentials, conductances, ...)
- `voltmeter` records the membrane potential
- `correlation_detector` records pairwise cross-correlations between the spiking activity of neurons
- `weight_recorder` records the weight of connections

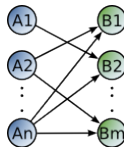
# SPECIFICATION OF CONNECTIVITY

- One-to-one
- All-to-all
- Fixed in-, out-degree
- Fixed total number ( $N$ )
- Pairwise Bernoulli ( $p$ )

```
A = Create('iaf_psc_alpha', n)  
B = Create('spike_detector', n)  
Connect(A, B, 'one_to_one')
```



```
A = Create("iaf_psc_alpha", n)  
B = Create("iaf_psc_alpha", m)  
Connect(A, B, {'rule': 'fixed_indegree',  
                'indegree': N})
```



# RANDOMIZATION OF SYNAPSE PROPERTIES

Dozens of standard distributions are natively supported.

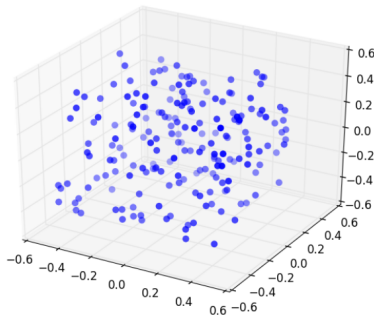
```
delay_dist = {'distribution': 'uniform',  
              'low': 0.8, 'high': 2.5}
```

```
alpha_dist = {'distribution': 'normal_clipped',  
              'low': 0.5, 'mu': 5.0,  
              'sigma': 1.0}
```


```
syn_dict = {'model': 'stdp_synapse',  
            'weight': 2.5,  
            'delay': delay_dist,  
            'alpha': alpha_dist}
```

# STRUCTURED NETWORKS USING TOPOLOGY

- Set node positions on grids or arbitrary points in space (1D, 2D, 3D)
- Nodes can be neurons or combinations of neurons and devices
- Connect nodes in a position- and distance-dependent manner
- Set boundary condition (periodic or not)



# GAP JUNCTIONS




Neuron  $i$  (hh\_psc\_alpha\_gap)

$y'_i(t) = f_i(y_i(t))$ ,  $y_i(t_0)$  given

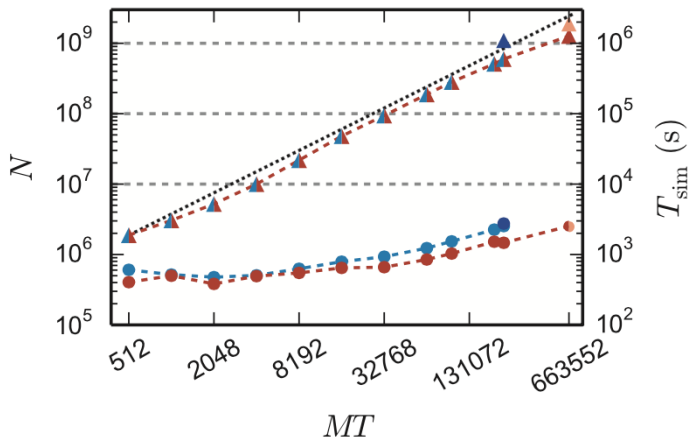
$$\begin{aligned} \frac{V'_i}{C_m} = & -I_i^{ionic}(V_i, m_i, h_i, n_i, p_i) \\ & + I_i^{applied}(I_i^{ex}, I_i^{in}) \\ & + I_i^{gap}(V_i, V_j) \end{aligned}$$

- at each time point neuron  $i$  needs membrane potential of neuron  $j$
- large system of differential equations
- naïve: communication of  $V$  in each step
- better: Jacobi waveform relaxation



Hahne et al. (2015). A unified framework for spiking and gap-junction interactions in distributed neural network simulations. *Frontiers in Neuroinformatics*. 9:22

# NEST PERFORMANCE



Maximum network size and corresponding run time as function of number of virtual processes on the K computer (red) and JUQUEEN (blue). Taken from Kunkel et al., (2014), [Front Neuroinf.](https://doi.org/10.3389/fninf.2014.00078) DOI: 10.3389/fninf.2014.00078

# HOW TO USE NEST?

- The Python interface PyNEST

```
import nest
```

```
n = nest.Create("iaf_psc_exp", 8000, {"V_m": -65.})
```

```
sd = nest.Create("spike_detector")
```

```
nest.Connect(n, sd)
```

```
...
```

```
nest.Simulate(1000.)
```

- The multi-simulator Python interface PyNN

```
from pyNN import nest
```

```
nest.setup()
```

```
neuron_t = nest.native_cell_type("iaf_psc_exp")
```

```
pop_exc = nest.Population(8000, neuron_t, {})
```



```
pop_exc.initialize(V_m=-65.)
```

```
...
```

```
nest.run(1000.)
```

# NEURONAL SIMULATIONS IN NEST

A simulation in NEST mimics a neuroscientific experiment.





# A FULL EXAMPLE

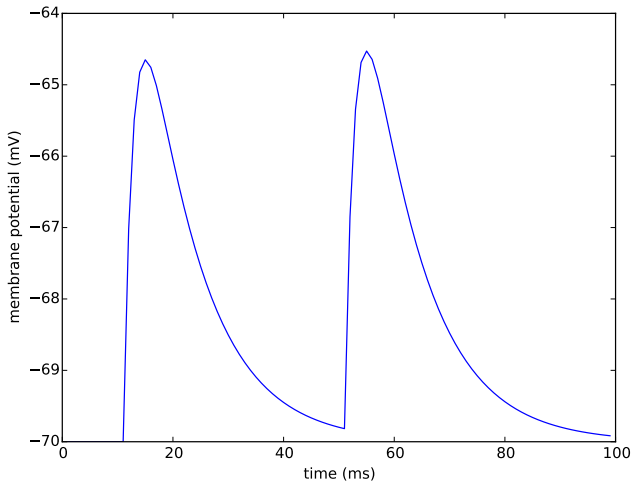
```
import nest # import NEST module
neuron = nest.Create('iaf_psc_exp') # create a neuron
voltmeter = nest.Create('voltmeter') # create a voltmeter
spikegenerator = nest.Create('spike_generator') # create a spike generator
nest.SetStatus(spikegenerator, {'spike_times': [10., 50.]}) # let it spike

# connect spike generator and voltmeter to the neuron
nest.Connect(spikegenerator, neuron, syn_spec={'weight' : 1E3})
nest.Connect(voltmeter, neuron)

nest.Simulate(100.) # run the simulation

# read out recording time and voltage from voltmeter and plot them
times = nest.GetStatus(voltmeter)[0]['events']['times']
voltage = nest.GetStatus(voltmeter)[0]['events']['V_m']
pl.plot(times, voltage)
pl.xlabel('time (ms)'); pl.ylabel('membrane potential (mV)')
pl.show()
```

# A FULL EXAMPLE



# EVENT-DRIVEN VS. TIME-DRIVEN


## Event-driven simulation:

- Visit a neuron only when it receives an event (e.g. a spike)
- From  $y(t_i)$ , calculate  $y(t_{i+1})$

## Time-driven simulation:

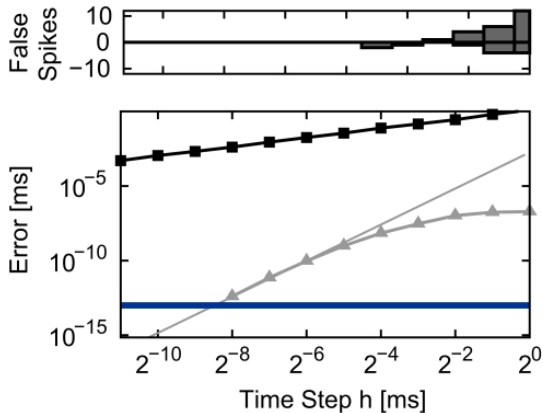
- Visit each neuron in each time step  $h$
- From  $y(ih)$ , calculate  $y([i + 1]h)$

# EVENT-DRIVEN VS. TIME-DRIVEN



Hanuschkin et al. Front. Neuroinf. 2010

# EVENT-DRIVEN VS. TIME-DRIVEN



Hanuschkin et al. Front. Neuroinf. 2010

# EVENT-DRIVEN VS. TIME-DRIVEN

	Event-driven	Time-driven
Pros	<ul style="list-style-type: none"><li>■ more efficient for low input rates</li><li>■ 'correct' solution for invertible neuron models</li></ul>	<ul style="list-style-type: none"><li>■ more efficient for high input rates</li><li>■ works for all neuron models</li><li>■ scales well</li></ul>
Cons	<ul style="list-style-type: none"><li>■ only works for neurons with invertible dynamics</li><li>■ event queue does not scale well</li></ul>	<ul style="list-style-type: none"><li>■ only 'approximate' solution even for analytically solvable models</li><li>■ spikes can be missed due to discrete sampling of membrane potential</li></ul>

# EVENT-DRIVEN VS. TIME-DRIVEN

NEST uses a hybrid approach to simulation

- input events to neurons are frequent: time-driven algorithm
  - If the dynamics is nonlinear, we need a numerical method to solve it, e.g.:
    - Forward Euler:  $y([i + 1]h) = y(ih) + h \cdot \dot{y}(ih)$
    - Runge-Kutta ( $k$ -th order)
    - Runge-Kutte-Fehlberg with adaptive step size
    - ...
- Use a pre-implemented solver, for example, from the GNU Scientific Library (GSL).
- If the dynamics is linear (e.g. leaky integrate-and-fire), we can solve it exactly.

# EVENT-DRIVEN VS. TIME-DRIVEN


NEST uses a hybrid approach to simulation

- input events to neurons are frequent: time-driven algorithm
  - If the dynamics is nonlinear, we need a numerical method to solve it, e.g.:
    - Forward Euler:  $y([i + 1]h) = y(ih) + h \cdot \dot{y}(ih)$
    - Runge-Kutta ( $k$ -th order)
    - Runge-Kutta-Fehlberg with adaptive step size
    - ...
- Use a pre-implemented solver, for example, from the GNU Scientific Library (GSL).
- If the dynamics is linear (e.g. leaky integrate-and-fire), we can solve it exactly.
- events at synapses are rare: event driven component
  - Exception: gap junctions



# MINIMUM SYNAPTIC DELAY


During an interval of the minimal transmission delay in the network ( $\Delta$ ), neurons are effectively decoupled.



Input events to neurons are frequent: time-driven updates.  
Events at synapses are rare: event driven.


# COMMUNICATION OF EVENTS

- communication only required in intervals of the minimal delay between neurons




# COMMUNICATION OF EVENTS

- communication only required in intervals of the minimal delay between neurons
- communication frequency independent of step size  $h$




# COMMUNICATION OF EVENTS

- communication only required in intervals of the minimal delay between neurons
- communication frequency independent of step size  $h$
- less communications containing more data is more efficient due to overhead of communication between machines



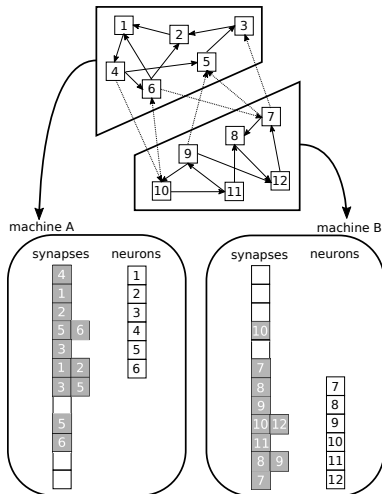
# REPRESENTATION OF NETWORK STRUCTURE: SERIAL

- Each neuron and synapse maintains its own parameters
- Synapses save the index of the target neuron



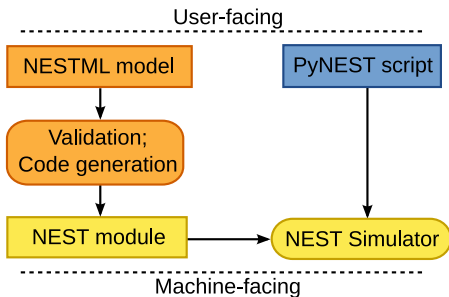
# REPRESENTATION OF NETWORK STRUCTURE: DISTRIBUTED

- neurons are distributed round robin onto processes
- one target list for every neuron on each machine
- synapse stored on machine that hosts the target neuron
- wiring is a parallel operation



# NESTML

NESTML is a domain-specific language for neuron and synapse models.



Using PyNEST, you instantiate and connect the models that you define in NESTML.

# NESTML: DESIGN PRINCIPLES

- Concise; low on boilerplate
- Speak in the vernacular of the neuroscientist (keywords such as neuron, synapse)
- Easy (dynamical) equation handling coupled with imperative-style programming (if  $V_m \geq \text{threshold}$ : ...)

NESTML comes with a code generation toolbox.

- Code generation (model definition but not instantiation)
- Automated ODE analysis and solver selection
- Flexible addition of targets using Jinja2 templates



# NESTML: EXAMPLE

```
neuron iaf_psc_exp:
```

```
  state:
```

```
    V_abs mV = 0 mV
```

```
  end
```

```
  equations:
```

```
    shape G = exp(-t / tau_syn)
```

```
    V_abs' = -V_abs / tau_m  
            + (I_ext + convolve(G, spikes)) / C_m
```

```
  end
```

```
  parameters:
```

```
    C_m      pF = 250 pF
```

```
    tau_m    ms = 10 ms
```

```
    tau_syn  ms = 2 ms
```

```
    V_threshold mV = 40 mV # w.r.t. zero!
```

```
  end
```

```
input:
```

```
  spikes pA <- spike
```

```
  I_ext pA <- current
```

```
end
```

```
update:
```

```
  integrate_odes()
```

```
  if V_abs > V_threshold:
```

```
    V_abs = 0
```

```
    emit_spike()
```

```
  end
```

```
end
```

```
end
```

# GETTING HELP

## Within Python:

```
nest.help('iaf_psc_exp')  
nest.help('Connect')
```

## Online documentation:

<https://nest-simulator.readthedocs.io/>

## Community:

- NEST user mailing list
- Bi-weekly open video conference
- <http://github.com/nest/nest-simulator/>
- Annual NEST Conference: a forum for users and developers

Please tell us about problems. We can only fix what we know of!

# REFERENCES AND FURTHER READING

- The NEST Simulator documentation at <https://nest-simulator.readthedocs.io/>
- Gewaltig et al. (2012) *NEST by example: An introduction to the neural simulation tool NEST*. doi:10.1007/978-94-007-3858-4\_18
- Hanuschkin et al. (2010) *A general and efficient method for incorporating precise spike times in globally time-driven simulations*. doi:10.3389/fninf.2010.00113
- Jordan et al. (2018) *Extremely Scalable Spiking Neuronal Network Simulation Code: From Laptops to Exascale Computers*. doi:10.3389/fninf.2018.00002

# ACKNOWLEDGMENTS

This presentation is based on previous work by many people.

- Hannah Bos
- David Dahmen
- Moritz Deger
- Jochen Martin Eppler
- Espen Hagen
- Abigail Morrison
- Jannis Schuecker
- Johanna Senk
- Tom Tetzlaff
- Sacha van Albada
- Charl Linssen